

Learning-Based Adaptive Navigation Controller for Unmanned Ground Vehicles

A Benchmark Study of Estimation, Planning, Control, and Learned Adaptation — Open-Source Research Layer

Abstract

We present a fully reproducible, open-source research pipeline for UGV navigation that benchmarks three state estimators (EKF, UKF, sliding-window factor graph), five global planners (A*, Theta*, Hybrid A*, RRT*, MPPI), and five path-tracking controllers (PID, Pure Pursuit, Stanley, NMPC, and a learning-adapted NMPC) on a common kinematic bicycle model with simulated noisy GPS/IMU/encoder sensing. A learned speed-scheduling policy, trained in 88 s of CPU time with the Cross-Entropy Method, reduces RMS cross-track error by 12% and steering-rate variance by 40% relative to the fixed-parameter NMPC, confirming the adaptation hypothesis that motivates a full PPO/SAC meta-policy. A closed-loop robustness study with the EKF in the control loop reveals that all high-precision controllers converge to a common error floor of about 0.065 m imposed by state-estimation accuracy, indicating that beyond Stanley-level precision, localization — not control — is the binding constraint. All code, seeds, and results are released.

1. Method Overview

Pipeline: sensor simulation (GPS 5 Hz sigma = 0.15 m; gyro with bias random walk; encoder noise) feeds the state-estimation layer; planners operate on 6 x 4 m occupancy grids (0.1 m resolution); controllers track a figure-eight reference (lemniscate, 4 x 2 m) on a kinematic bicycle ($L = 0.32$ m) with rate-level inputs $[a, d\delta]$ shared by ALL controllers for fair smoothness comparison; the learning layer adapts NMPC parameters online. The NMPC solves a 30-step, 50 ms multiple-shooting OCP (CasADi/ipopt) with $Q = \text{diag}(10, 10, 5, 1, 0.1)$, $R = \text{diag}(0.5, 2, 0)$, delta-u penalty, and warm starting — the identical OCP later compiled with ACADOS SQP-RTI for ROS2 deployment.

2. State Estimation Benchmark

Estimator	RMS_pos_m	std	RMS_psi_rad	Mean_ms
EKF	0.0876	0.0044	0.0355	0.0280
UKF	0.0874	0.0043	0.0354	0.1073
FactorGraph	0.1101	0.0106	0.0646	9.1130

Table 1: 5 seeds, 40 s runs. UKF matches EKF once the sigma-point mean uses a circular mean for heading; the smoother trails in this GPS-rich scenario.




Figure 1: position error over time and RMS summary with seed variance.

3. Planner Benchmark

Planner	Time_ms	Length_m	Smoothness_rad	Success_ %
A*	3.27	6.50	11.78	100.00
Theta*	64.15	6.17	1.11	100.00
Hybrid A*	11.47	6.10	0.83	100.00
RRT*	1154.46	6.28	1.34	100.00
MPPI	3362.85	6.28	7.40	100.00

Table 2: 3 maps x 3 seeds, 100% success for all planners. Hybrid A* is smoothest and kinematically feasible at 18 ms.




Figure 2: planned paths on map 1; planning time (log) and length/smoothness bars.

4. Controller Benchmark with Learned Adaptation

Method	RMS_ect_m	RMS_jerk	SteerVar	MeanSpeed	Progress_m	p99_ms
PID	0.2262	0.1024	0.1929	1.4784	58.5750	0.0303
PurePursuit	0.1879	0.1024	0.0489	1.4784	58.5750	0.1402
Stanley	0.0269	0.1024	0.0537	1.4784	58.5750	0.0384
MPC	0.0187	0.7451	0.0234	1.7758	70.4206	8.3312
Adaptive-MPC	0.0165	0.8436	0.0139	1.3761	54.6313	8.5141

Table 3: full 40 s figure-eight. Adaptive-MPC: CEM-learned curvature speed schedule $v_{ref} = v_{base} / (1 + k_{curv} |\kappa|)$, converged to $v_{base} = 1.89$, $k_{curv} = 1.35$.

Learning-Based Adaptive Navigation — Full Controller Benchmark




Figure 3: trajectories, cross-track error, velocity (with learned schedule), and metric bars.

5. Robustness: Estimation in the Control Loop

Method	GT_state_RMS_m	EKF_state_RMS_m	std	Degradation_x
PurePursuit	0.1879	0.1574	0.0407	0.8378
Stanley	0.0269	0.0655	0.0029	2.4331
MPC	0.0187	0.0657	0.0108	3.5065
Adaptive-MPC	0.0165	0.0637	0.0062	3.8634

Table 4: true tracking error when each controller acts on the EKF estimate (3 noise seeds) vs ground-truth state.




Figure 4: precision controllers collapse to a common ~ 0.065 m error floor set by estimation accuracy; Pure Pursuit improves because its lookahead low-pass filters estimation noise.

6. Discussion

Three findings. (i) Adaptation works: even a 2-parameter learned schedule beats hand-tuned NMPC on error and steering smoothness, trading average speed — the trade-off a full PPO/SAC meta-policy will manage contextually. (ii) The estimation floor: below ~ 0.07 m of estimation error, controller precision is masked; investment should go to localization (or estimator-aware control, motivating solver/estimator-health signals in the RL observation). (iii) Benchmark integrity is fragile: we found and fixed three silent bugs — reference lobe-jumping at the path self-intersection, linear angle averaging in the UKF sigma-point mean, and closed-path index clamping that stopped the robot after one lap while INFLATING its error score. Each would have corrupted the conclusions without visible failure.

7. Limitations and Next Steps

Kinematic model without tire dynamics; sensor simulation in place of Gazebo/Isaac physics; CEM policy has 2 parameters vs the specified 58-dim PPO/SAC meta-policy; ipopt (10-14 ms) vs the ACADOS RTI target (< 5 ms). Next: Gazebo sensor plugins replace simulation/sensors.py, PPO/SAC training (rl/ppo_training.py, environment smoke-tested) on GPU, ACADOS port with the < 2 cm trajectory-match validation gate, and the 2,400-episode ROS2 benchmark matrix.

Reproducibility: fixed seeds throughout; tests/test_all_modules.py verifies all six modules; results CSVs committed alongside plots.